

# GitOps for Apache Kafka® At scale and with ease using Jikkou!

A dark, atmospheric illustration of a futuristic landscape. In the foreground, there are tall, spiky grasses. In the middle ground, a valley leads towards a city skyline featuring many tall, illuminated skyscrapers. The sky is a mix of deep blues and warm orange and yellow hues from the setting sun. There are a few small, distant figures or vehicles on the ground.

The Opensource Resource as Code Framework for Apache Kafka®



@fhussonnois

# Florian Hussonnois

Lead Software Engineer @Kestra.

- Worked for 12 years as Consultant & Trainer
- 10+ years of experience with Apache Kafka (version 0.7)
- Open-source projects such as Kafka Connect File Pulse, Jikkou
- Confluent Community Catalyst since 2019



   @fhussonnois

# The True Story Behind Every Kafka Resources!

It all starts with Kafka's CLI.

```
#!/bin/bash
kafka-topics --bootstrap-server localhost:9092 \
  --create --topic my_topic \
  --partitions 6 \
  --replication-factor 1

# (output)
# Created topic my-topic.
```

*in a galaxy far, far away...*

**Padawan:** "Can you please create a new topic for my team?"

**Jedi:** "Yes, of course! How many partitions do you need?"

**Padawan:** "...(silence)"





# What About Configuration ?

Yet another story!

```
#!/bin/bash
# Update to new value
kafka-configs --bootstrap-server localhost:9092 \
--entity-type topics --entity-name my_topic \
--alter --add-config retention.ms=17280000
```

```
#!/bin/bash
# Reset to default value
kafka-configs --bootstrap-server localhost:9092 \
--entity-type topics --entity-name my_topic \
--alter --delete-config retention.ms
```

*in a galaxy far, far away...*

**Amiral:** "Could you update the config of this topic?"

**Stormtrooper:** "Yes, give me a second, please."

**Amiral:** "Why did we lose all our data?"

**Stormtrooper:** "Oops, we forgot a zero in retention.ms."

# What's the problem ?

With Manual Configuration Management

## No versioning

- How to keep track of changes over time?

## Non-repeatable

- How to reproduce modifications on multiple environments ?

## No traceability

- Who did what and when ?

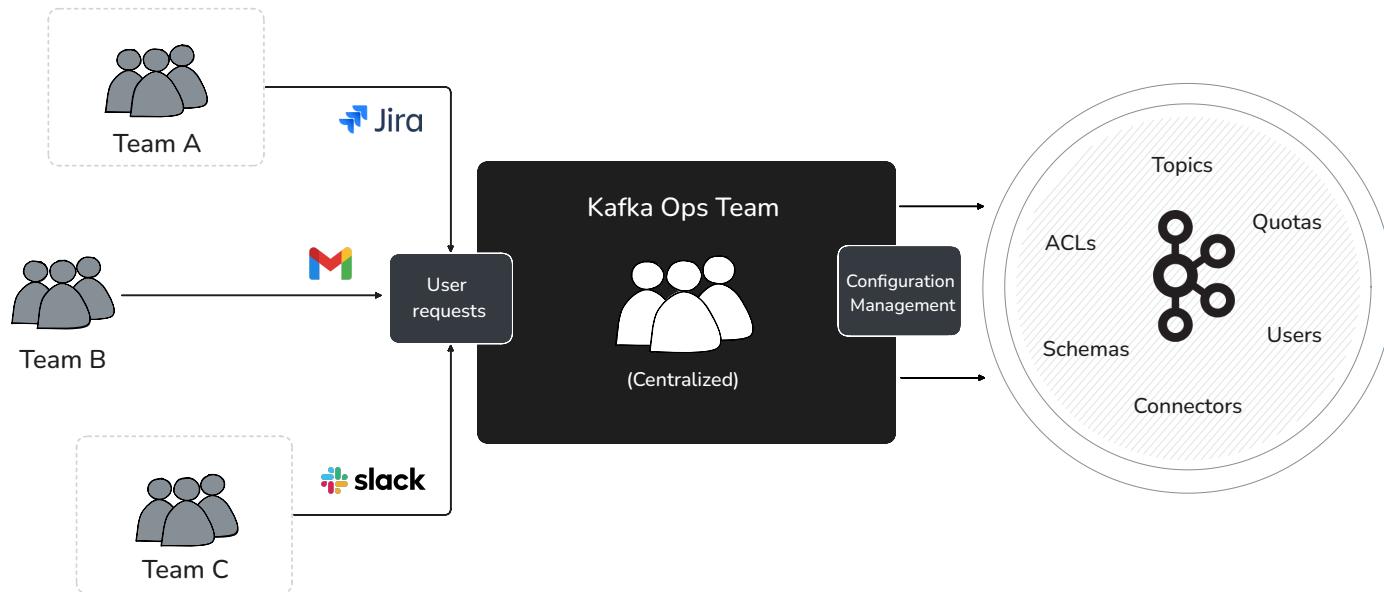
## No human fault-tolerance

- One typo, and boom!



# Hard to Scale!

Centralized team can quickly become a bottleneck



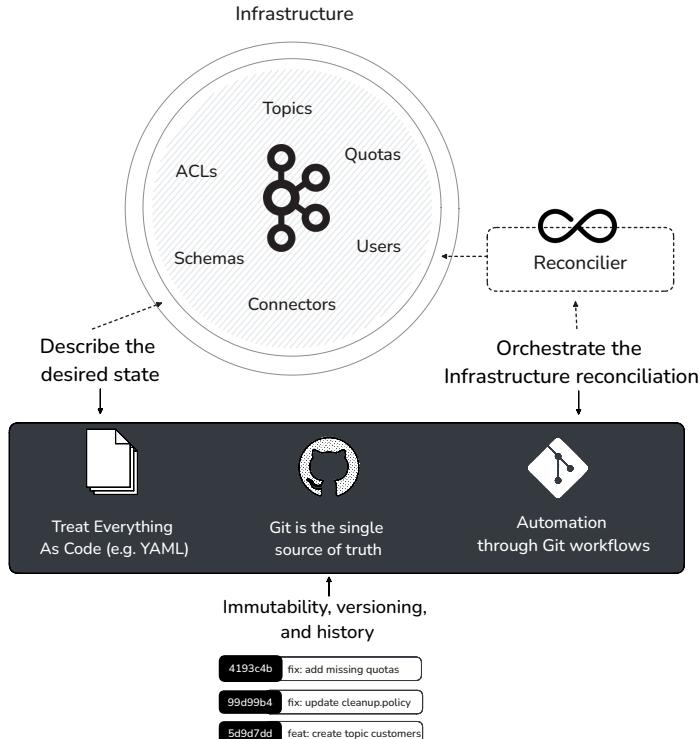
- Creating new topics or requesting access takes hours instead of minutes.

# The GitOps and IaC Approach

Everything As Code!

**Synchronize** the desired states (*plan*) of our resources and configurations with the actual state of our infrastructure from Git.

- Enables a decentralized approach - multiple repositories for each team.
- Shifts responsibilities to development teams.
- Higher autonomy



# Existing IaC Solutions

Why they don't fit ?

- Terraform (or OpenTofu)
  - Often preferred by **DevOps** teams, but not as commonly used by development teams.
  - Have to learn HCL - Multiple providers for Kafka
- Strimzi (Kubernetes Operators) - Helm
  - Not everyone runs on Kubernetes, or has the ability to deploy an operator.

Not specialized solutions: no control over the pushed configurations.



# Introducing Jikkou

The Opensource Resource as Code Framework for Apache Kafka®



# What is Jikkou ?

Jikkou is a flexible framework enabling developers and Devops teams to efficiently manage, automate and provision all the resources required for their projects.

- Open Source - **Apache License, Version 2.0.**
- Written in Java.
- Declarative (YAML) - Resource as Code.
- Easy to versioning, Tracking changes over time.
- GitOps - Continuously Reconciliation
- Extensible - Developed for **Apache Kafka®**, but designed for anything!



*Jikkou (jikkō/実行), means "execution (e.g. of a plan), carrying out, (putting into) practice in Japanese.*

# Everything As A Resource

Express the desired state of resources using YAML descriptor files.

```
---  
# ./my-topic.yaml  
apiVersion: "kafka.jikkou.io/v1beta2"  
kind: KafkaTopic  
metadata:  
  name: 'my-topic'  
spec:  
  partitions: 3  
  replicas: 1  
  configs:  
    min.insync.replicas: 1  
    cleanup.policy: 'delete'
```

- Same resource model as Kubernetes to describe the entities to manage.



# Jikkou original goal

A modern and intuitive command line client for Apache Kafka.

Jikkou was initially developed as a simple **Command Line Interface (CLI)** to efficiently **Create, Read, Update , Delete**, and version resources such as Topics, ACLs, and Quotas.

```
# Installing Jikkou CLI using SDKMAN!
$ sdk install jikkou
```

Available as a **native image** (built with GraalVM) or **Java Binary distribution**.

# Jikkou CLI

## CORE COMMANDS:

|          |  |
|----------|--|
| apply    | Update the resources as described by the resource definition files.                                |
| create   | Create resources from the resource definition files (only non-existing resources will be created). |
| delete   | Delete resources that are no longer described by the resource definition files.                    |
| diff     | Show resource changes required by the current resource definitions.                                |
| get      | Display one or many specific resources.  |
| prepare  | Prepare the resource definition files for validation.  |
| update   | Create or update resources from the resource definition files                                      |
| validate | Check whether the resources definitions meet all validation requirements.                          |

## SYSTEM MANAGEMENT COMMANDS:

|        |                                      |
|--------|--------------------------------------|
| action | List/execute actions.                |
| health | Print or describe health indicators. |

## ADDITIONAL COMMANDS:

|                     |  |
|---------------------|--|
| api-extensions      | Print the supported API extensions                   |
| api-resources       | Print the supported API resources                    |
| config              | Sets or retrieves the configuration of this client   |
| generate-completion | Generate bash/zsh completion script for jikkou.      |
| help                | Display help information about the specified command |

# Configuration

HOCON

```
# ./application.conf
jikkou {
    extension.providers {
        default.enabled: true
    }

    kafka {
        client {
            bootstrap.servers = "localhost:9092"
            bootstrap.servers = ${?KAFKA_BOOTSTRAP_SERVERS}  $ jikkou config use-context localhost
        }
    }
}
```

## Configure context

```
$ jikkou config set-context localhost \
--config-file "`pwd`/application.conf"
```

## Use context

## Display current configuration

```
$ jikkou config view
```

See: <https://github.com/lightbend/config>

# Reconcile

Apply resource changes

## Multiple reconciliation modes:

- `CREATE` : Create new resource.
- `UPDATE` : Create or update existing resources.
- `DELETE` : Delete existing resource.
- `FULL` : Apply any changes.

NOTE: Jikkou supports a `dry-run` mode.

## To only display state changes

```
jikkou diff \  
--files ./my-topic.yaml \  
--selector "metadata.labels.environment IN (demo)" \  
--output YAML
```

## To only create new resources

```
jikkou create \  
--files ./my-topic.yaml \  
--selector "metadata.labels.environment IN (demo)" \  
--output YAML
```

# Apply

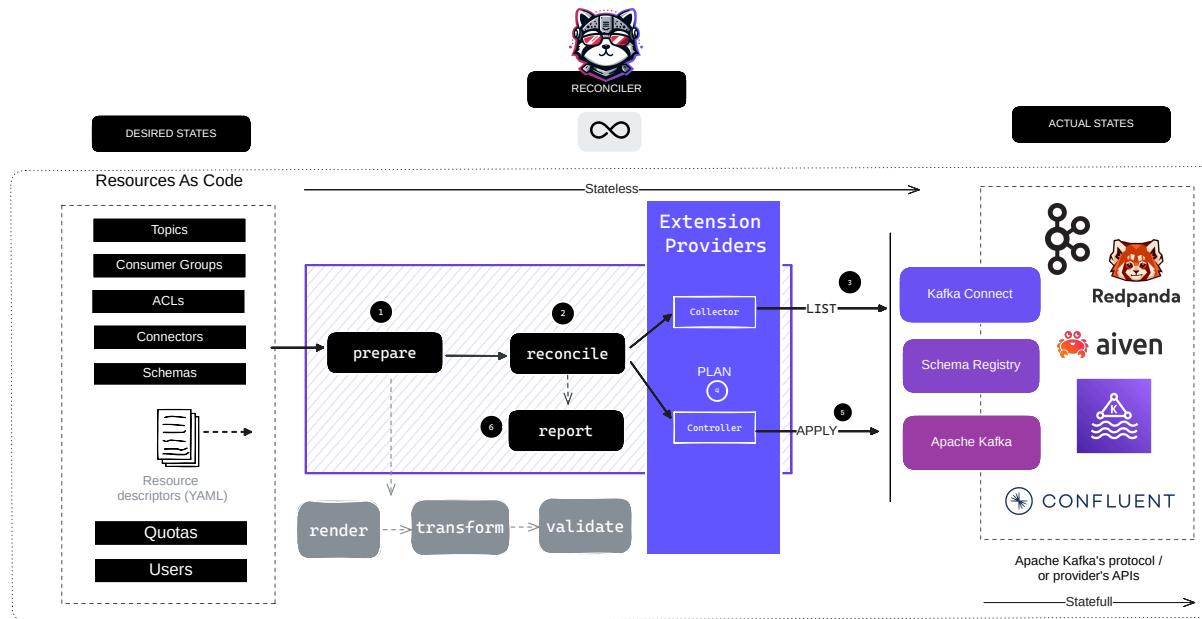
Resource State Changes

(output)

```
---  
kind: "ApiChangeResultList"  
apiVersion: "core.jikkou.io/v1"  
metadata:  
  labels: {}  
  annotations:  
    jikkou.io/changes-count: 1  
dryRun: false  
results:  
- end: "2023-12-20T00:00:00.000000Z"  
  status: "CHANGED"  
  description: "Create topic 'my-topic' (partitions=3, replicas=1, configs=[cleanup.policy=delete,min.insync.replicas=1])"  
  changed: true  
  failed: false  
  change:  
    apiVersion: "kafka.jikkou.io/v1beta2"  
    kind: "KafkaTopicChange"  
    metadata:  
      name: "my-topic"
```

# How It Works ?

Leverage your system as the Source of Truth!



Jikkou adopts a stateless approach and does not store any state internally.

- Seamless integration with other solutions.

# Templating

Jikkou provides a simple templating mechanism based-on Jinjava, a [Jinja template engine for Java](#).

## Template Resource File

```
# ./kafka-topics.tpl
---
apiVersion: 'kafka.jikkou.io/v1beta2'
kind: 'KafkaTopicList'
items:
{% for country in values.countryCodes %}
- metadata:
    name: "{{ values.topicPrefix}}-iot-events-{{ country }}"
  spec:
    partitions: {{ values.topicConfigs.partitions }}
    replicas: {{ values.topicConfigs.replicas }}
    config:
      retention.ms: 3600000
      max.message.bytes: 20971520
{% endfor %}
```

## Data Values File

```
# ./values.yaml
---
topicConfigs:
  partitions: 4
  replicas: 1
topicPrefix: "{{ system.env.TOPIC_PREFIX | default('test', true) }}"
countryCodes:
  - fr
  - uk
  - it
```

## Command

```
#!/bin/bash
jikkou prepare \
--files kafka-topics-template.tpl \
--values-files kafka-topics-values.yaml
```

# Validate

Ensure that inbound resources conform to specific rules or constraints.

**Example:** Validate that topic names match a specific regex.

```
# ./application.conf
jikkou {
    validations = [
        {
            name = "topicMustHaveValidName"
            type = io.streamthoughts.jikkou.kafka.validation.TopicNameRegexValidation
            priority = 100
            config = {
                topicNameRegex = "[a-zA-Z0-9\\.\\_\\-]+"
            }
        }
    ]
}
```

# Transform

Transform, enrich, or filter inbound resources

**Example:** Enforce a minimum value for the replication factor of kafka topics.

```
# ./application.conf
jikkou {
    transformations: [
        {
            type = io.streamthoughts.jikkou.kafka.transform.KafkaTopicMinReplicasTransformation
            priority = 100
            config = {
                minReplicationFactor = 3
            }
        }
    ]
}
```

# Actions

Allow a user to execute a specific and one-shot operation on resources.

## Using Jikkou CLI:

```
$ jikkou action <ACTION_NAME> execute [<options>]
```

## Built-in actions:

- `KafkaConsumerGroupsResetOffsets`
  - Reset offsets of consumer group to earliest, lastest, a specific offset or from datetime.
- `KafkaConnectRestartConnectors`
  - Restarts all or just the failed Connector and Task instances for one or multiple named connectors.

"Ok, it's very cool! But...

How to use Jikkou if you don't have a direct public access to the your Kafka cluster?"

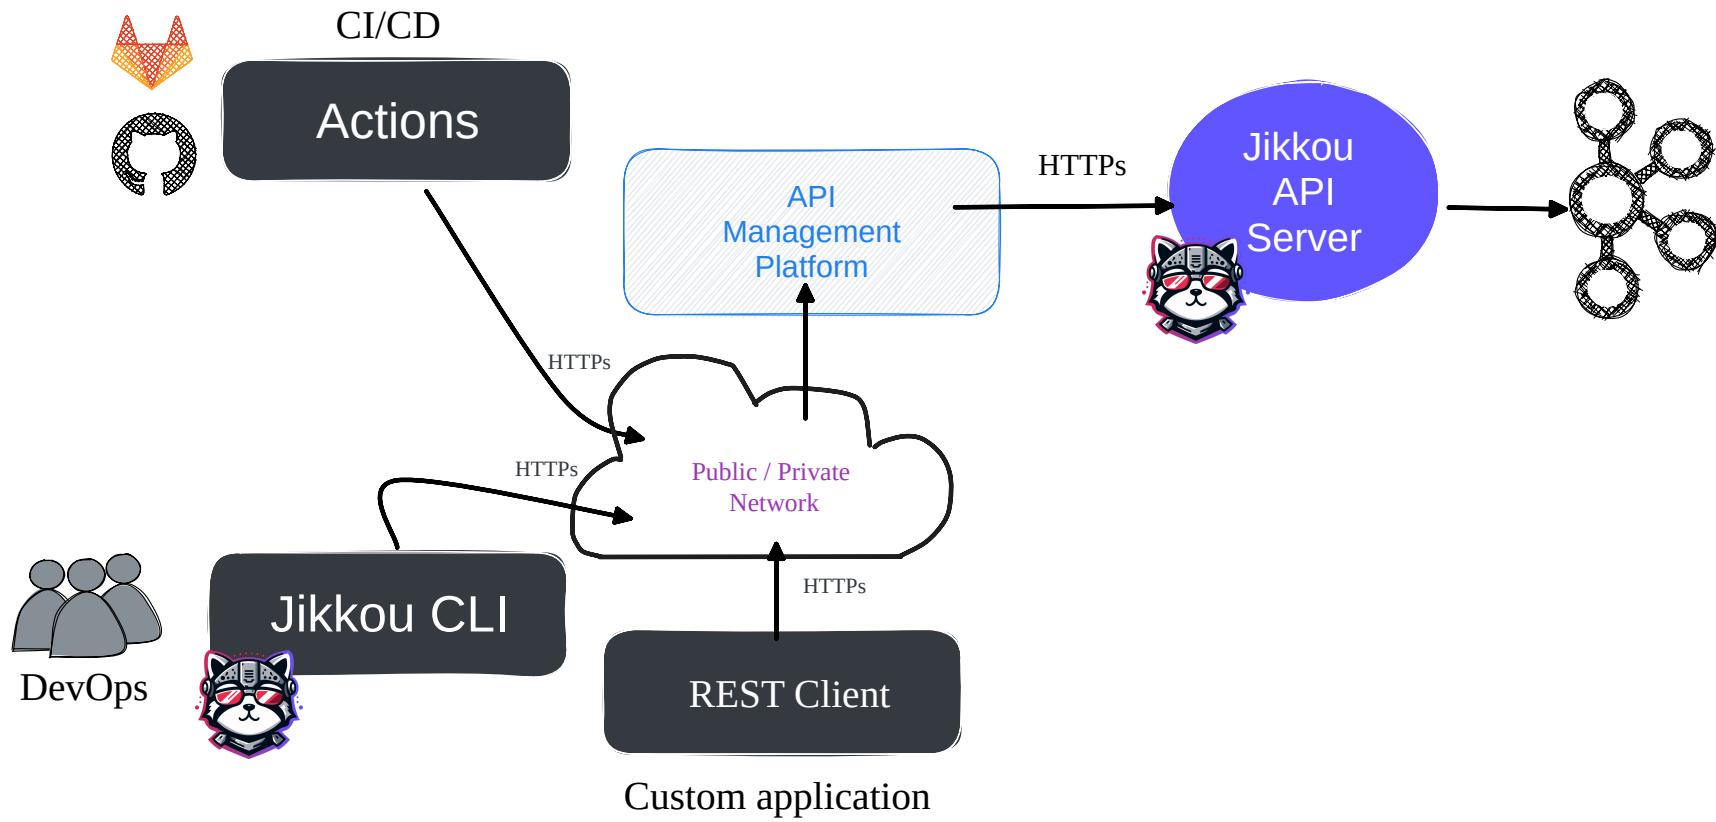
# Jikkou API Server

A REST interface that makes it easy to manage and automate all your data platform resources.

```
$ docker run -it --net host streamthoughts/jikkou-api-server:latest
```

# Jikkou API Server

Usage



# Jikkou API Server

## Key Benefits

- **Improved Security :**

- Can sit behind your API management platform.
- Provides additional security and auditing capabilities.

- **Centralized Governance :**

- Acts as a gateway to your data platform.
- Allows to centralize validation/transformation rules to manage resources.

- **Can be extended :**

- Supports custom resources and extensions.



# REST APIs

Jikkou API Server is built with Micronaut Framework

## **Listing resources.**

```
`GET /apis/{group}/{version}/{plural}`
```

## **Creating resources.**

```
`POST /apis/{group}/{version}/{plural}/reconcile	mode/{mode}{?dry-run}`
```

# Security

- Micronaut Framework supports multiple authentication mechanisms.
- Jikkou API Server with Basic Authentication, JWT, X.509 Certificate, OAuth2, etc.



# Jikkou CLI / Proxy Mode

Do not change your developer experience

```
jikkou {  
    # Proxy Configuration  
    proxy {  
        # Specify whether proxy mode is enabled (default)  
        enabled = true  
        # URL of the API Server  
        url = "http://localhost:28082"  
    }  
}
```

# DEMO

# Learn More

Documentations ·  GitHub ·  Blog Posts

# Thank You

Questions ?

★ Star the project on GitHub: <https://github.com/streamthoughts/jikkou>

